

Spark Streaming



cerenode.io

Overview

- Used for processing real-time streaming data
- Enables high-throughput and fault-tolerant stream processing of live data streams
- Fundamental unit is DStreams, which is basically a series of RDDs to process the real-time data



Steps



Streaming Context

- Consumes a stream of data in Spark
- Registers an InputDStream to produce a Receiver object
- Main entry for Spark Streaming functionality
- Provides number of default implementations of sources like Twitter, Akka Actor and ZeroMQ that are accessible from the context

Streaming Context Initialization

- Can be created from a SparkContext object

```
import org.apache.spark._  
import org.apache.spark.streaming._  
new StreamingContext(sc, Seconds(1))
```

Streaming Context Initialization

- Can be created from a SparkConf object

```
import org.apache.spark._  
import org.apache.spark.streaming._  
  
val conf = new  
SparkConf().setAppName(appName).setMaster(master)  
  
val ssc = new StreamingContext(conf, Seconds(1))
```




DStream

- Discretized Stream (DStream) is the basic abstraction provided by Spark Streaming
- It is a continuous stream of data
- It is received from source or a processed data stream generated by transforming the input stream
- Internally, a DStream is represented by a continuous series of RDDs and each RDD contains data from a certain interval



DStream



Input DStreams and Receivers

- Input DStreams represents stream of input data received from streaming sources
- Two categories sources
 - Basic: Sources directly available in the StreamingContext
 - Advanced: Sources like Kafka, Flume, Kinesis, etc.
- Every input DStream (except file stream) is associated with a Receiver object which receives the data from a source and stores it in Spark's memory for processing.

Basic Sources

- **SocketTextStream**
 - `StreamingContext.socketTextStream("<hostname>", <port>)`
- **File Streams**
 - `streamingContext.fileStream[KeyClass, ValueClass, InputFormatClass](dataDirectory)`
 - `streamingContext.textFileStream(dataDirectory)`


Transformations on DStream

- Allows the data from input DStream to be modified similar to RDDs
- Supports many of the normal transformations available for normal RDDs
- Common transformations
 - `map`, `flatMap`, `filter`, `reduce`, `groupBy`, `foreachRDD`, `updateStateByKey`



DStream Window

- Spark Streaming provides windowed computations which allow us to apply transformations over a sliding window of data



DStream Window

Some of the common window operations are as follows..

- `window(windowLength, slideInterval)`
- `countByWindow(windowLength, slideInterval)`
- `reduceByWindow(func, windowLength, slideInterval)`
- `reduceByKeyAndWindow(func, windowLength, slideInterval, [numTasks])`



Output Operations on DStream

- Output operations allow DStream's data to be pushed out to external systems like databases or file systems
- Output operations trigger the actual execution of all the DStream transformations
- Common output operations are
 - `print()`
 - `saveAsTextFiles(prefix, [suffix])`
 - `saveAsObjectFiles(prefix, [suffix])`
 - `foreachRDD(func)`

updateStateByKey

The `updateStateByKey` operation allows you to maintain arbitrary state while continuously updating it with new information. To use this, you will have to do two steps.

1. Define the state - The state can be an arbitrary data type.
2. Define the state update function - Specify with a function how to update the state using the previous state and the new values from an input stream.



Checkpointing

Spark Streaming needs to *checkpoint* enough information to a fault- tolerant storage system such that it can recover from failures.

There are two types of data that are checkpointed.

- *Metadata checkpointing* - Saving of the information defining the streaming computation to fault-tolerant storage like HDFS. This is used to recover from failure of the node running the driver of the streaming application (discussed in detail later). Metadata includes:
 - *Configuration* - The configuration that was used to create the streaming application.
 - *DStream operations* - The set of DStream operations that define the streaming application.
 - *Incomplete batches* - Batches whose jobs are queued but have not completed yet.
- *Data checkpointing* - Saving of the generated RDDs to reliable storage.